SELinux doesn’t bite

How to write SELinux policy
for your project painlessly

I lied, SELinux bites!

And the door is locked now,
you have to to stay.

Lukáš Zapletal

@lzap

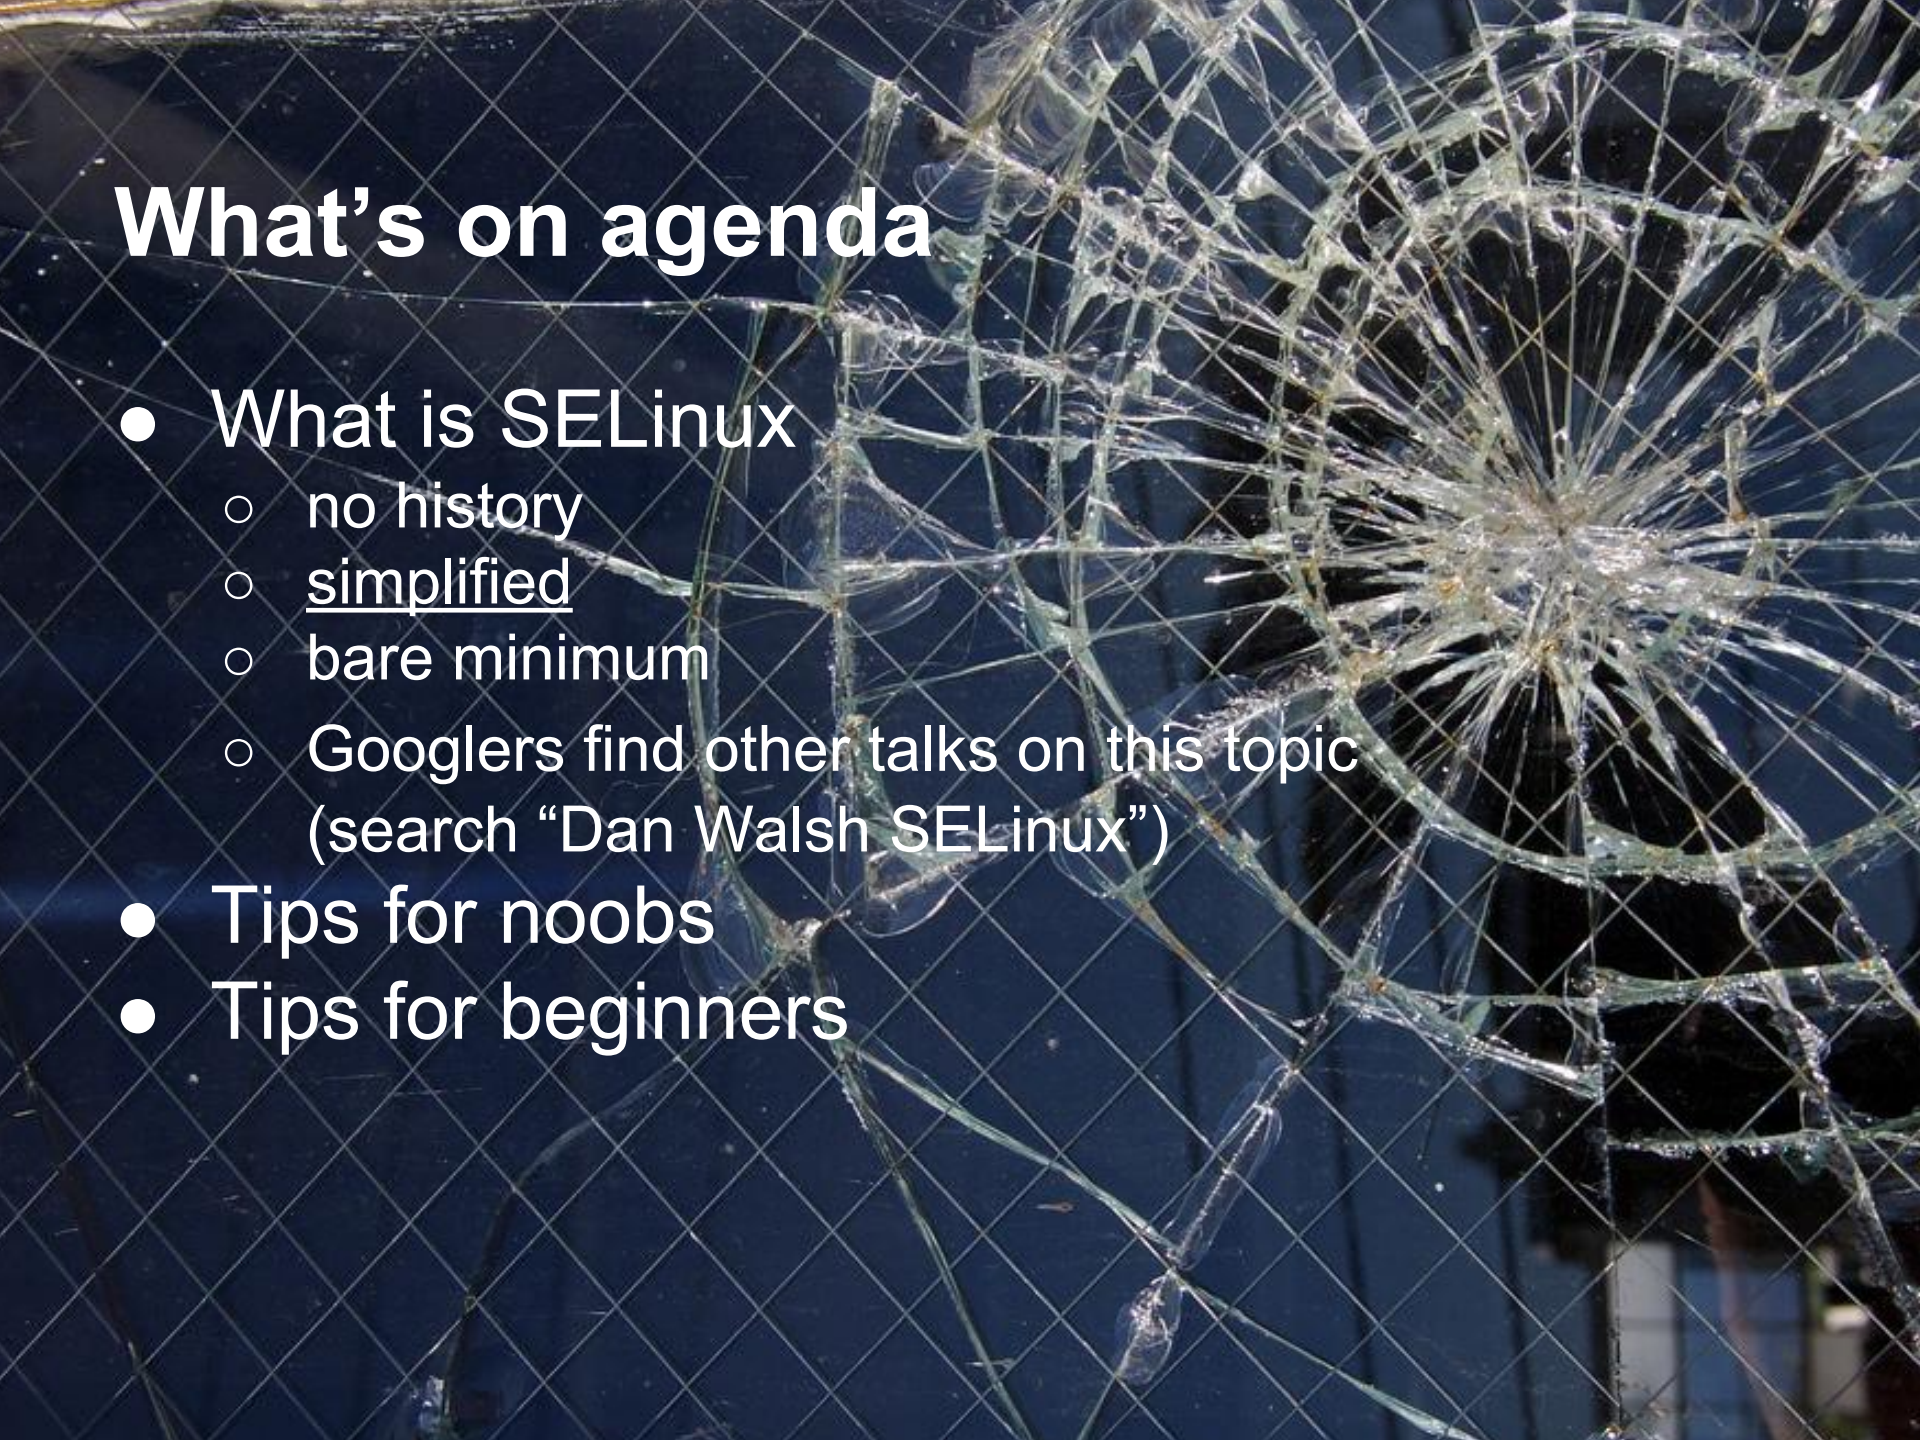

What's on agenda

- What is SELinux
  - no history
  - simplified
  - bare minimum
  - Googlers find other talks on this topic (search "Dan Walsh SELinux")
- Tips for noobs
- Tips for beginners

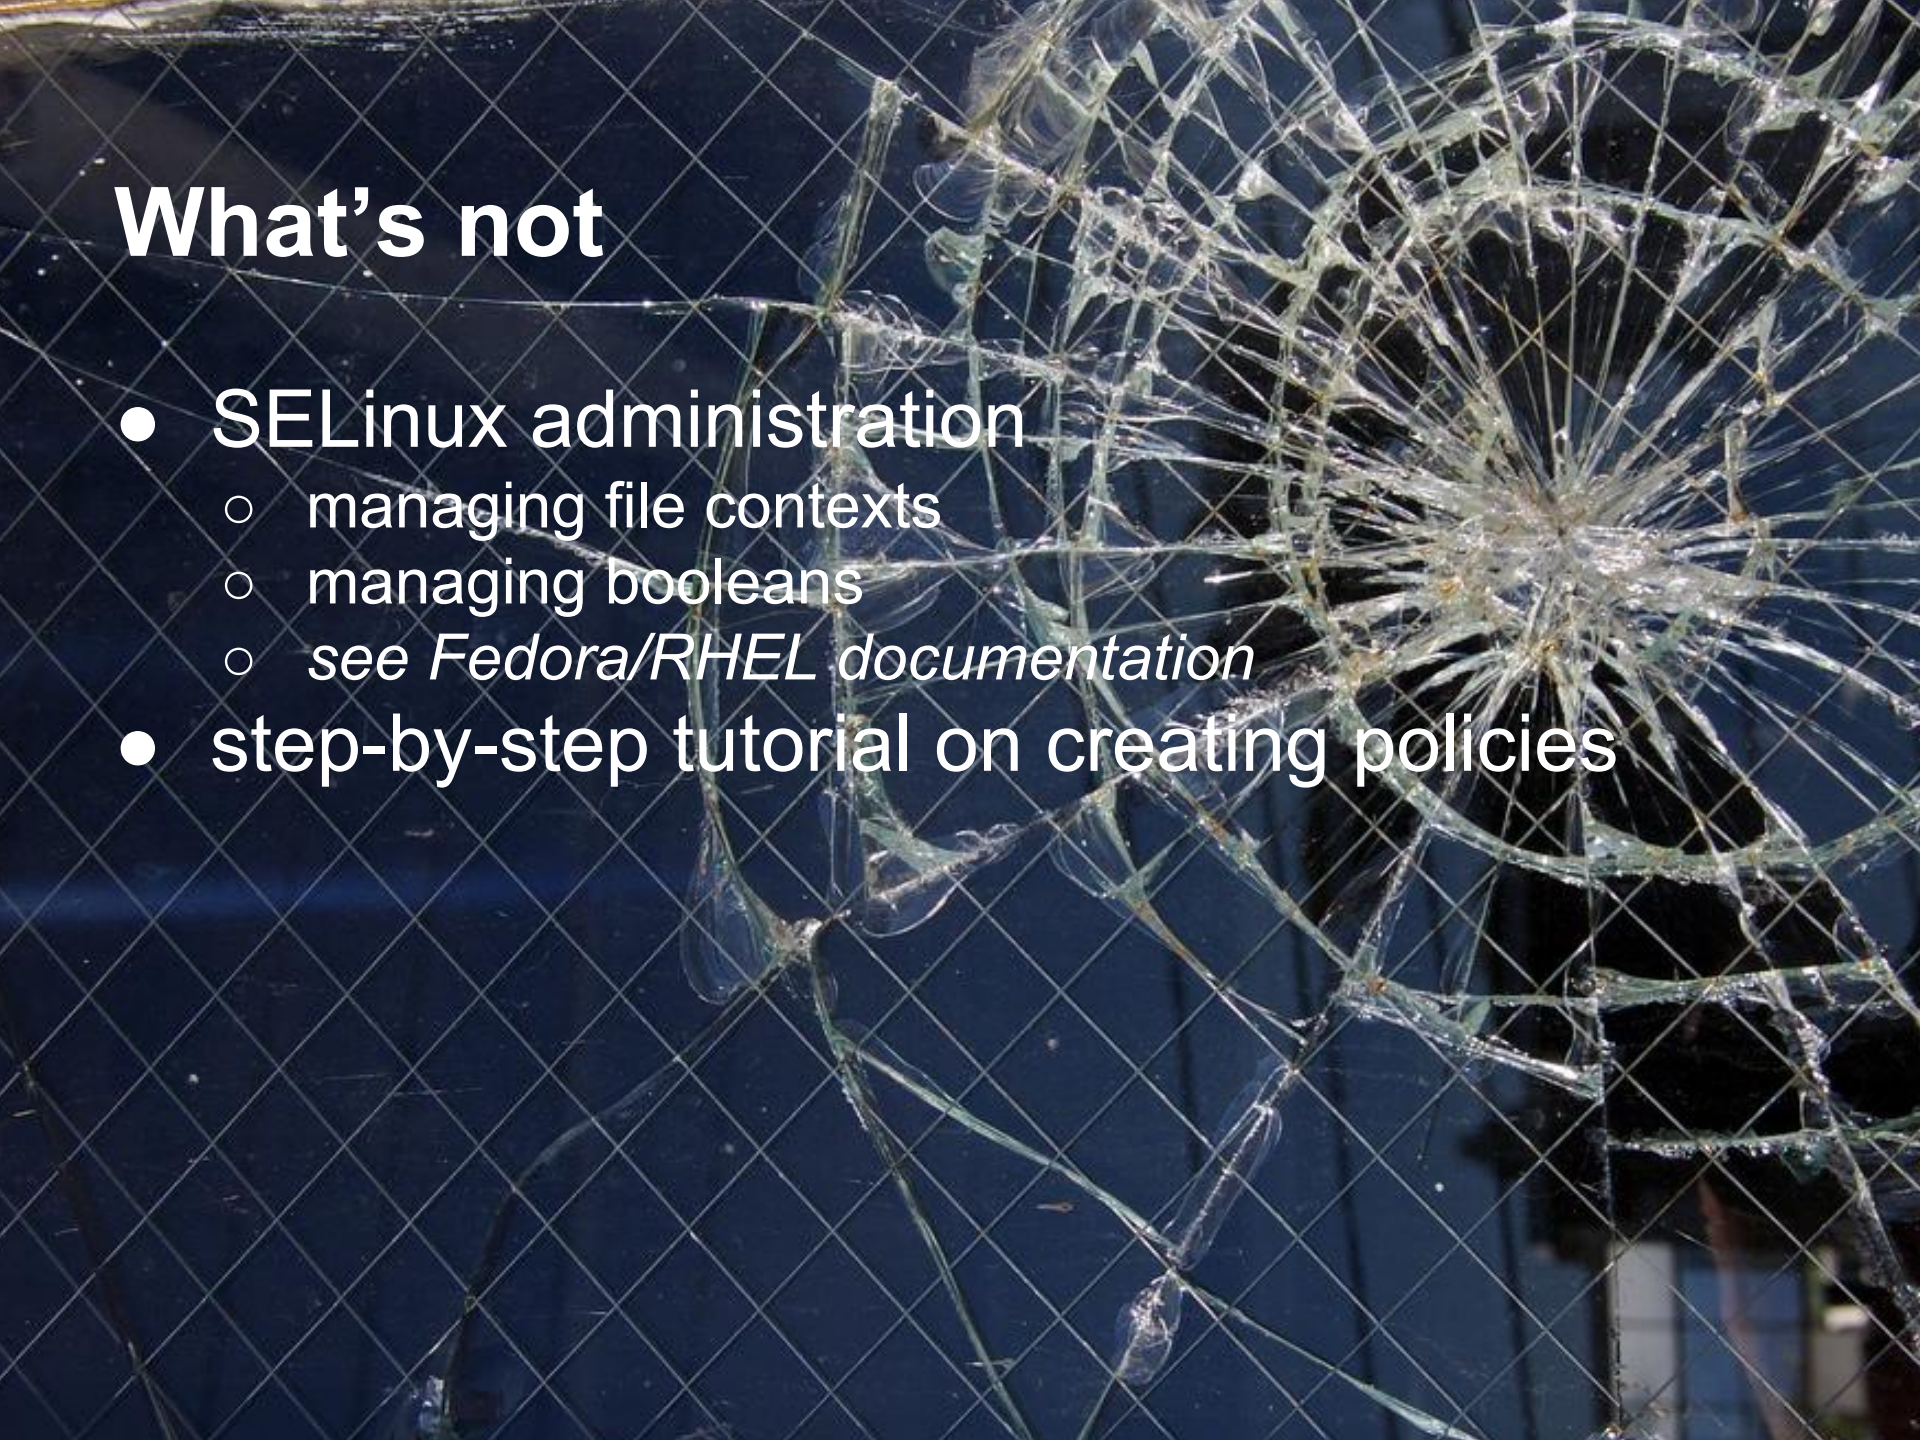

What's not

- SELinux administration
  - managing file contexts
  - managing booleans
  - *see Fedora/RHEL documentation*
- step-by-step tutorial on creating policies

What is SELinux

Linux kernel **module** that **enforces** mandatory access-control **policies**.

What is SELinux

SELinux makes sure that **subject** (process) does **follow** granular **set of rules**.

What is SELinux

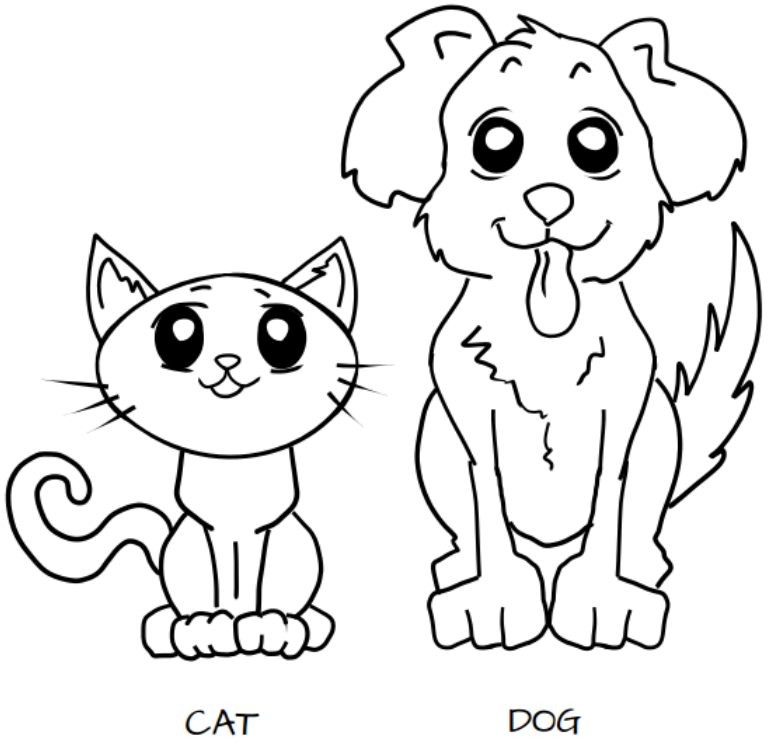

What is SELinux

What is SELinux

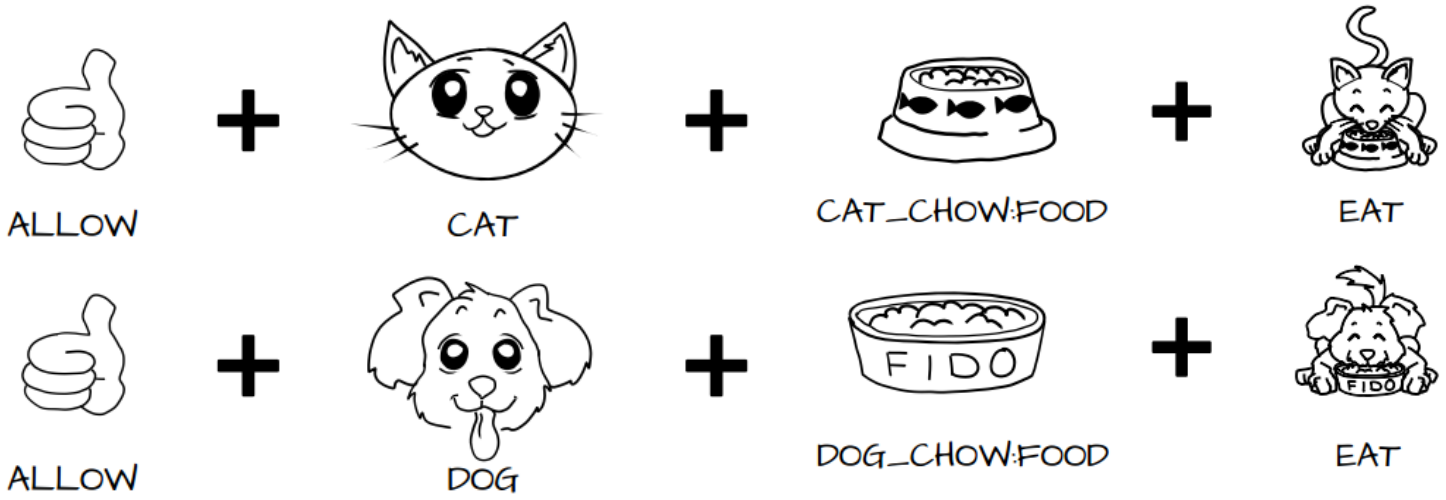

What is SELinux

What is SELinux

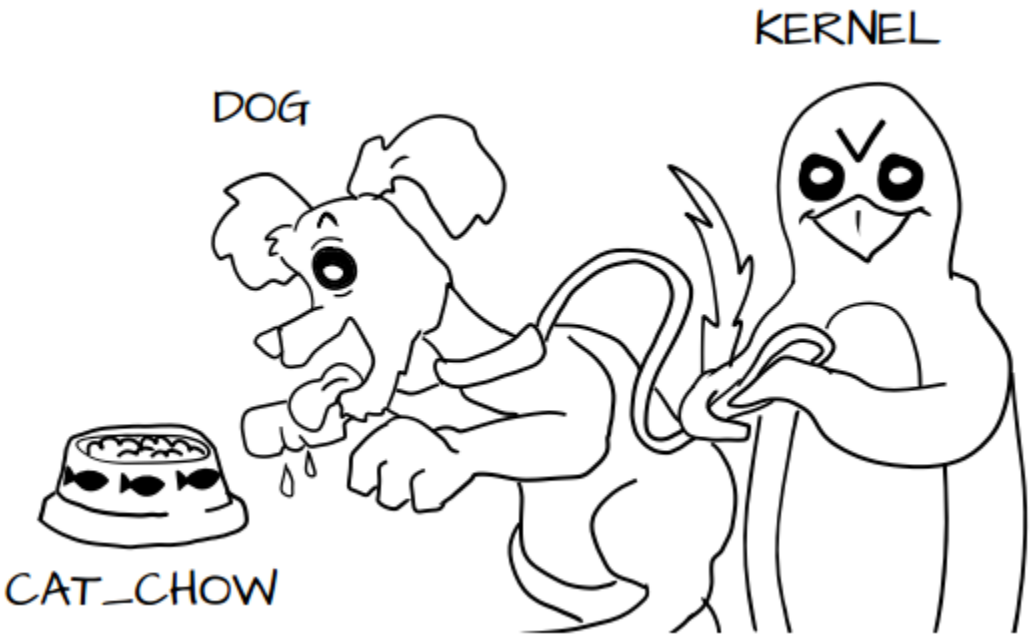

What is SELinux

MCS (Multi Category Security)

MLS (Multi Level Security)

What can SELinux do for you

- increases security
  - prevents from attacks (sql injection vs shellshock)
  - restricts investigations after successful attacks (open remote port)
  - warns during attacks (denials)

What can SELinux do for you

- find software bugs
  - unchecked file open return values
  - leaked descriptors
- workarounds proprietary behavior

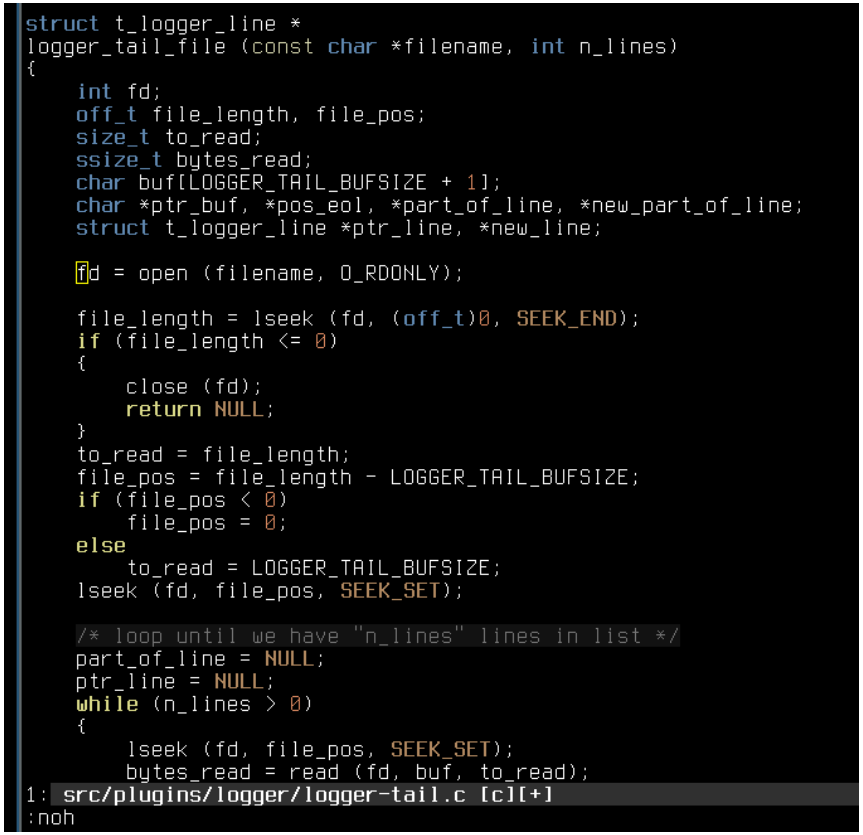

```
struct t_logger_line *
logger_tail_file (const char *filename, int n_lines)
{
    int fd;
    off_t file_length, file_pos;
    size_t to_read;
    ssize_t bytes_read;
    char buf[LOGGER_TAIL_BUFSIZE + 1];
    char *ptr_buf, *pos_eol, *part_of_line, *new_part_of_line;
    struct t_logger_line *ptr_line, *new_line;

    fd = open (filename, O_RDONLY);

    file_length = lseek (fd, (off_t)0, SEEK_END);
    if (file_length <= 0)
    {
        close (fd);
        return NULL;
    }
    to_read = file_length;
    file_pos = file_length - LOGGER_TAIL_BUFSIZE;
    if (file_pos < 0)
        file_pos = 0;
    else
        to_read = LOGGER_TAIL_BUFSIZE;
    lseek (fd, file_pos, SEEK_SET);

    /* loop until we have "n_lines" lines in list */
    part_of_line = NULL;
    ptr_line = NULL;
    while (n_lines > 0)
    {
        lseek (fd, file_pos, SEEK_SET);
        bytes_read = read (fd, buf, to_read);
1: src/plugins/logger/logger-tail.c [c][+]
:noh
```

SELinux policy in Fedora

SELinux policy in Fedora

```
$ rpm -qa selinux-policy*
selinux-policy-3.12.1-196.fc20.noarch
selinux-policy-targeted-3.12.1-196.fc20.noarch
selinux-policy-devel-3.12.1-196.fc20.noarch


$ rpm -ql selinux-policy-targeted
…
/etc/selinux/targeted/contexts/files/file_contexts
…
/etc/selinux/targeted/modules/active/modules/abrt.pp
/etc/selinux/targeted/modules/active/modules/apache.pp
…
```

SELinux policy in Fedora

```
$ rpm -ql selinux-policy-devel
…
/usr/share/man/man8/sshd_selinux.8.gz
…
/usr/share/selinux/devel/html/telnetd.html
…
/usr/share/selinux/devel/Makefile
/usr/share/selinux/devel/include/Makefile
…
/usr/share/selinux/devel/include/contrib/postfix.if
/usr/share/selinux/devel/include/kernel/corecommands.if
/usr/share/selinux/devel/include/system/iptables.if
…
/usr/share/selinux/devel/include/support/ipc_patterns.spt
…
```

SELinux custom policy - hello world

- mypolicy.te (type enforcement)
- mypolicy.if (interfaces and docs)
- mypolicy.fc (file contexts)

```
# touch mypolicy.{te,if,fc}
# echo “policy_module(mypolicy, 0.1)” > *te
# make -f /usr/share/selinux/devel/Makefile
# semodule -i mypolicy.pp
# semodule -l | grep mypolicy
mypolicy 0.1
```

SELinux custom policy - makefile

Default makefile targets

- all (compile, generate docs, load)
- load/reload
- refresh (reload all policies)
- clean

Important variables:

- NAME (targeted, minimum, mls)
- TYPE (standard, mls, mcs)
- QUIET (set to “n” for verbose output)

Example SELinux policy

myapp.te:

```
policy_module(myapp, 1.0.0)

# Declarations

type myapp_t;
type myapp_exec_t;
domain_type(myapp_t)
domain_entry_file(myapp_t, myapp_exec_t)

type myapp_log_t;
logging_log_file(myapp_log_t)

type myapp_tmp_t;
files_tmp_file(myapp_tmp_t)

# Myapp local policy

allow myapp_t myapp_log_t:file {  read_file_perms append_file_perms };

allow myapp_t myapp_tmp_t:file  manage_file_perms;
files_tmp_filetrans(myapp_t, myapp_tmp_t, file)
```

Example SELinux policy

myapp.if:

```
interface(`myapp_domtrans',`
  gen_require(`
    type myapp_t, myapp_exec_t;
  ')

  domtrans_pattern($1, myapp_exec_t, myapp_t)
')

interface(`myapp_read_log',`
  gen_require(`
    type myapp_log_t;
  ')

  logging_search_logs($1)
  allow $1 myapp_log_t:file  read_file_perms;
')
```

Example SELinux policy

myapp.fc:

```
/usr/sbin/myapp    --  gen_context(system_u:object_r:myapp_exec_t,s0)
/var/log/myapp     -d  gen_context(system_u:object_r:myapp_log_t,s0)
```

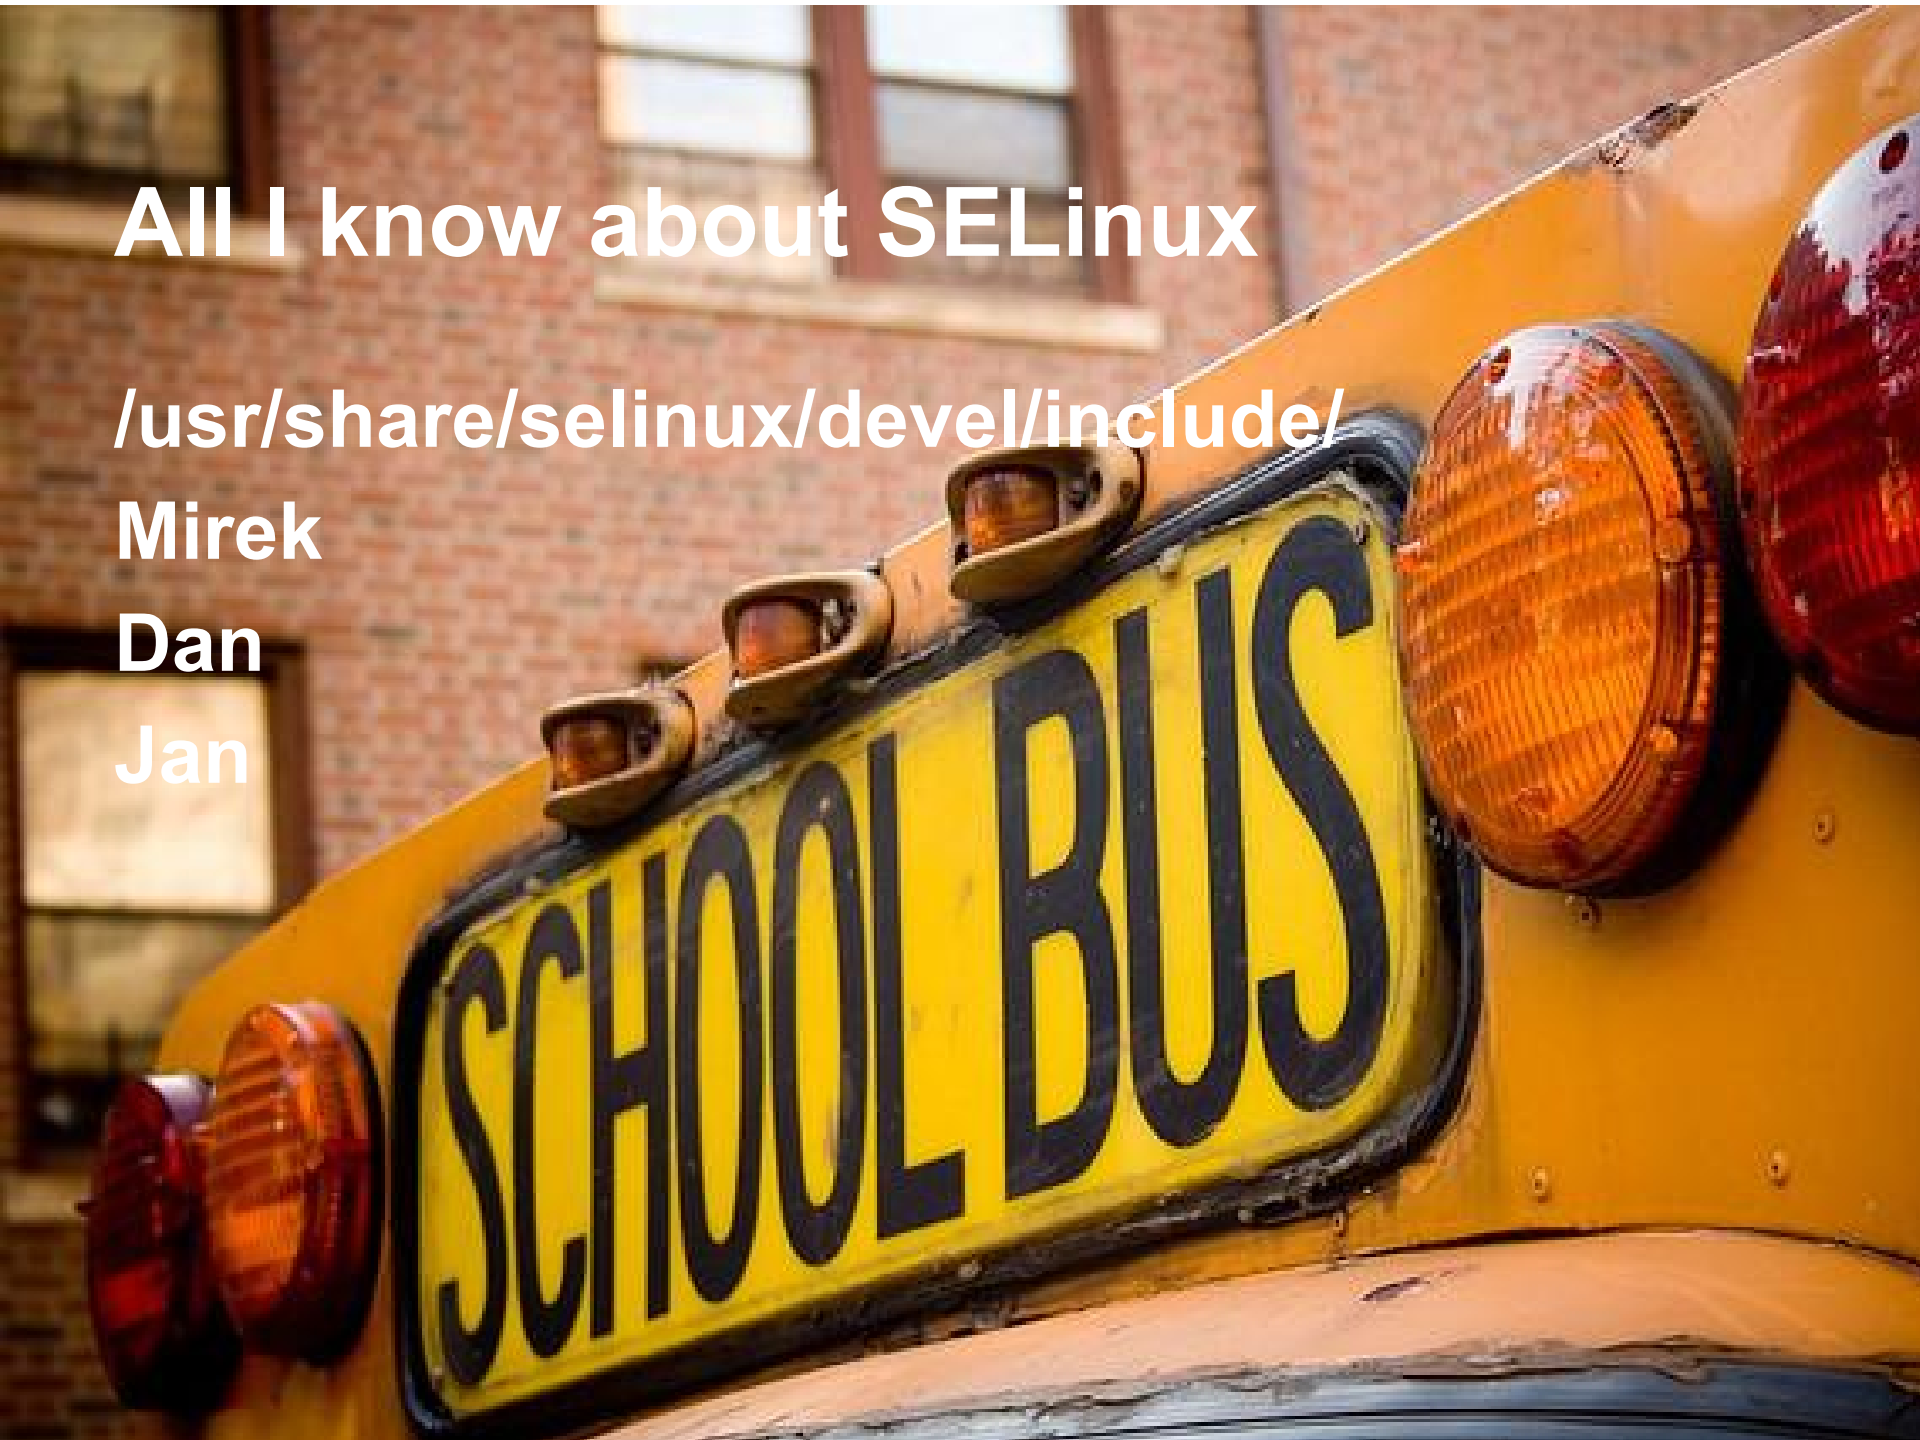

All I know about SELinux

/usr/share/selinux/devel/include/

Mirek

Dan

Jan

Important interface files

- application.if
- corenetwork.if
- files.if
- miscfiles.if
- devices.if
- terminal.if

- apache.if
- abrt.if

Important support files

- file_patterns.spt
- misc_macros.spt
- misc_patterns.spt
- loadable_module.spt

```
$ find /usr/share/selinux/devel/include -name \*.if | wc -
l
474
$ find /usr/share/selinux/devel/include -name \*.spt | wc
-l
8
```

This m4 preprocessor

This m4 preprocessor

<lzap> I was never big fan of m4 you know
<mgrepl> there’s upstream effort to replace it
<lzap> \o/
<mgrepl> it’s LISP-based language
<lzap> ***censored***
* lzap has left the channel (weechat 0.4.1)

This m4 preprocessor

```
/usr/bin/checkmodule:  loading policy configuration from tmp/foreman.tmp
foreman.te":238:ERROR 'syntax error' at token 'xxx_pattern'  on line 10522:
    xxx_pattern(passenger_t, httpd_tmp_t, httpd_tmp_t)
```

This m4 preprocessor

```
#line 238

#line 238
                } # end require
#line 238

#line 238

#line 238
        if (httpd_run_foreman) {
#line 238

#line 238
    manasge_dirs_pattern(passenger_t, httpd_tmp_t, httpd_tmp_t)
#line 238

#line 238
        allow passenger_t httpd_tmp_t:dir { open read getattr lock search ioctl add_name remove_name write };
#line 238
        allow passenger_t httpd_tmp_t:file { create open getattr setattr read write rename link unlink ioctl lock };
#line 238

#line 238

#line 238
        allow passenger_t httpd_tmp_t:dir { open read getattr lock search ioctl add_name remove_name write };
#line 238
        allow passenger_t httpd_tmp_t:sock_file { create open getattr setattr rename link unlink ioctl lock append };
#line 238
```

When to semicolon with m4

```
allow blah_t blahblah_t:file { read };
```

VS

```
myapp_read_blahblah_files(blah_t)
```

Interface naming

```
# from files.if
interface(`files_read_usr_files',`
   gen_require(`
      type usr_t;
   ')

   allow $1 usr_t:dir list_dir_perms;
   read_files_pattern($1, usr_t, usr_t)
   read_lnk_files_pattern($1, usr_t, usr_t)
')
```

Interface naming

```
# from obj_perm_sets.spt
define(`list_dir_perms',
`{ getattr search open read lock ioctl }')
```

Interface naming

```
# from file_patterns.spt
define(`read_files_pattern',`
   allow $1 $2:dir search_dir_perms;
   allow $1 $3:file read_file_perms;
')

define(`read_lnk_files_pattern',`
   allow $1 $2:dir search_dir_perms;
   allow $1 $3:lnk_file read_lnk_file_perms;
')
```

Searching for interface definitions

- code examples
- free download
- functions.sh
  - seshowif
  - sefindif
  - seshowdef
  - sefinddef

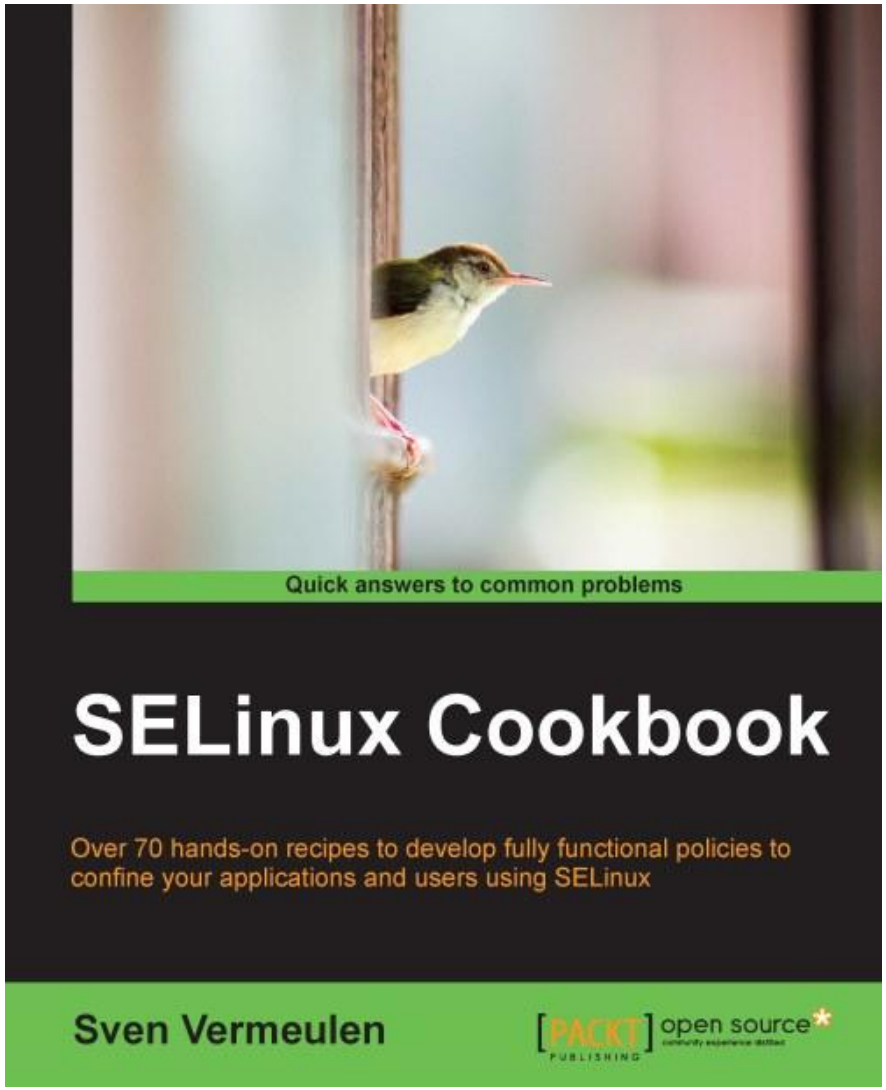

Searching for interface definitions

```
$ seshowif logging_log_file
interface(`logging_log_file',`
        gen_require(`
                attribute logfile;
        ')

        files_type($1)
        files_associate_tmp($1)
        fs_associate_tmpfs($1)
        typeattribute $1 logfile;
')

$ seshowdef search_dir_perms
define(`search_dir_perms',`{ getattr search open }')
```

Searching for interface definitions

$ sefindif logging_log_file

…

```
contrib/pki.if: template(`pki_apache_template',`
contrib/pki.if:         logging_log_file($1_log_t)
contrib/pki.if:         logging_log_filetrans($1_t, $1_log_t, { file dir } )
contrib/razor.if: template(`razor_common_domain_template',`
contrib/razor.if:       logging_log_filetrans($1_t, razor_log_t, file)
contrib/sendmail.if: interface(`sendmail_create_log',`
contrib/sendmail.if:    logging_log_filetrans($1, sendmail_log_t, file)
contrib/tomcat.if: template(`tomcat_domain_template',`
contrib/tomcat.if:      logging_log_file($1_log_t)
contrib/tomcat.if:      logging_log_filetrans($1_t, $1_log_t, { dir file })
kernel/files.if: interface(`files_stub_tmp',`
kernel/files.if: ##              <li>logging_log_file()</li>
system/authlogin.if: interface(`auth_log_filetrans_login_records',`
system/authlogin.if:    logging_log_filetrans($1, wtmp_t, file)
system/logging.if:
system/logging.if: ##           <li>logging_log_filetrans()</li>
system/logging.if: ##   logging_log_file(mylogfile_t)
system/logging.if: ##   logging_log_filetrans(mydomain_t, mylogfile_t, file)
system/logging.if: interface(`logging_log_file',`
```

…

How to navigate through with ctags

```
#!/bin/bash

/bin/rpm -q ctags > /dev/null

if [ $? == 0 ]; then
    if [ -d /usr/share/selinux/devel ]; then
        ctags -e --langdef=te --langmap=te:..te.if.spt \
            --regex-te='/^type[ \t]+(\w+)(,|;)/\1/t,type/' \
            --regex-te='/^typealias[ \t]+\w+[ \t+]+alias[ \t]+(\w+);/\1/t,type/' \
            --regex-te='/^attribute[ \t]+(\w+);/\1/a,attribute/' \
            --regex-te='/^[ \t]*define\(`(\w+)/\1/d,define/' \
            --regex-te='/^[ \t]*interface\(`(\w+)/\1/i,interface/' \
            --regex-te='/^[ \t]*bool[ \t]+(\w+)/\1/b,bool/'\
            /usr/share/selinux/devel/include/*/*.if \
            /usr/share/selinux/devel/include/support/*.spt *.te
    else
        echo "You need to install selinux-policy-devel package"
        exit 1
    fi
else
    echo "You need to install ctags package"
    exit 1
fi
```

You lucky Vim user!

https://github.com/lzap/vim-selinux

Anatomy of SELinux denial

 # grep AVC /var/log/audit/audit.log

type=AVC msg=audit(1413987601.193:1489): avc: denied { **name_bind** } for pid=12828 comm="ruby" src=**1251** scontext=system_u:system_r:**passenger_t**:s0 tcontext=system_u:object_r:**unreserved_port_t**:s0 tclass=**udp_socket**

 # ausearch -m AVC

--

type=AVC msg=audit(1413987601.193:1489): avc: denied { name_bind } for pid=12828 comm="ruby" src=1251 scontext=system_u:system_r:passenger_t:s0 tcontext=system_u:object_r:unreserved_port_t:s0 tclass=udp_socket

type=SYSCALL msg=audit(1413987601.193:1489): arch=x86_64 syscall=bind success=no exit=EACCES a0=b a1=**7f5438524080** a2=10 a3=0 items=0 ppid=1 pid=12828 auid=4294967295 uid=997 gid=995 euid=997 suid=997 fsuid=997 egid=995 sgid=995 fsgid=995 tty=(none) ses=4294967295 comm=ruby exe=/opt/rh/ruby193/root/usr/bin/ruby subj=system_u:system_r:passenger_t:s0 key=(null)

--

The audit2allow thing

```
# audit2allow -al
allow passenger_t unreserved_port_t:udp_socket name_bind;

# audit2allow -Ral
corenet_udp_bind_generic_port(passenger_t)

# audit2allow -R
<paste> Ctrl+D

# audit2allow -RalM quickfix
******************** IMPORTANT ***********************
To make this policy package active, execute:

semodule -i quickfix.pp
```

The audit2allow abuse

permissive + audit2allow =

The audit2allow abuse

- file contexts
- domain transitions
- software bugs are hidden
- not following the least privilege principle

SELinux policy artifacts

- the policy itself
- the process
  - design issues
  - misconfigurations
  - bugs

Take small steps

- Modify
- Compile
- Load
- Commit
- Repeat

One commit one issue (w/ denial)

```
commit 2a8011b2d211a043868c1bf3cff3d0dd084575eb
Refs: [docker-port-8989]
Author:     Lukas Zapletal <lzap+git@redhat.com>
AuthorDate: Fri Jan 16 10:34:44 2015 +0100
Commit:     Lukas Zapletal <lzap+git@redhat.com>
CommitDate: Fri Jan 16 10:34:44 2015 +0100

    Fixes #8989 - Add docker_port_t port and boolean

    Boolean passenger_can_connect_docker allows connections to newly created
    docker_port_t which is not yet defined in RHEL7/Fedora. This can be used when
    users starts Docker on TCP (defaults to UNIX sockets). Ports were reserved
    by IANA 2015-01-09: http (2375), https (2376).

    Denial:

    type=AVC msg=audit(1421352630.245:15331): avc:  denied  { name_connect } for
    pid=4803 comm="ruby" dest=2375 scontext=unconfined_u:system_r:passenger_t:s0
    tcontext=system_u:object_r:port_t:s0 tclass=tcp_socket
```

Who should write policies?

Review code (at least in two)

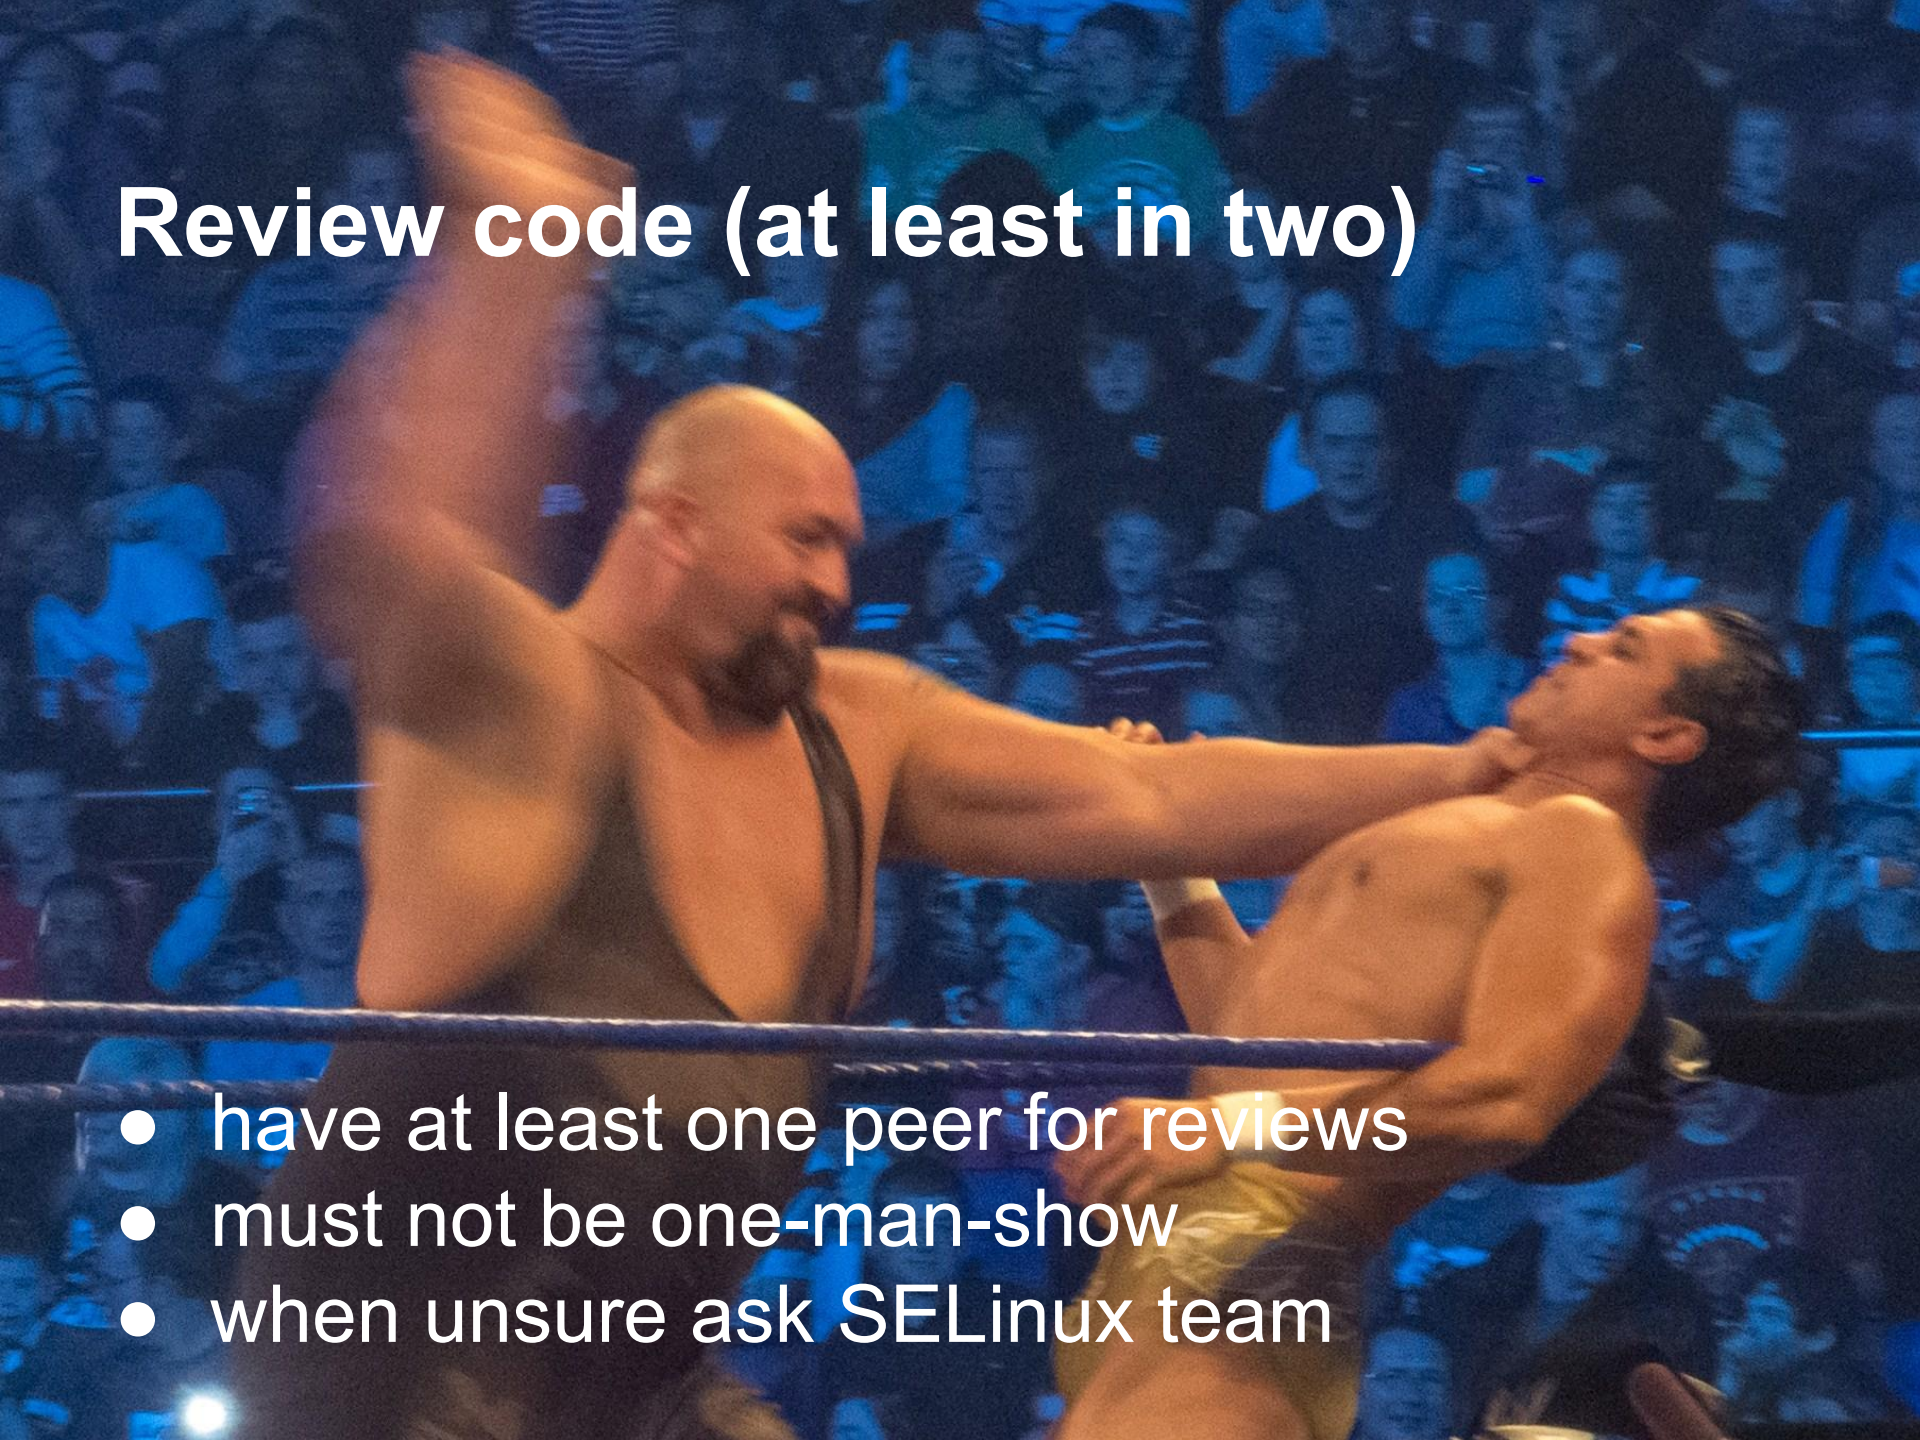

- have at least one peer for reviews
- must not be one-man-show
- when unsure ask SELinux team

Multiple distributions

```
# tcp connect to default OpenStack keystone API (5000)
ifdef(`distro_rhel6', `
  corenet_tcp_connect_commplex_port(passenger_t)
',`
  corenet_tcp_connect_commplex_main_port(passenger_t)
')


# cat Makefile
make -C ${TMPDIR} \
-f /usr/share/selinux/devel/Makefile \
DISTRO=rhel7
```

Deployment

[lzap@lzapx foreman-selinux]$ cat foreman-selinux-enable

```
#!/bin/bash
set +e

TMP=$(mktemp -t foreman-selinux-enable.XXXXXXXXXX)
trap "rm -rf '$TMP'" EXIT INT TERM

selinuxvariant=targeted


if /usr/sbin/semodule -s $selinuxvariant -l >/dev/null; then

  /usr/sbin/semanage module -S $selinuxvariant \
    -a /usr/share/selinux/${selinuxvariant}/foreman.pp.bz2

  echo "boolean -m --on httpd_setrlimit" > $TMP

  /usr/sbin/semanage port -E | grep -q elasticsearch_port_t || \
    echo "port -a -t elasticsearch_port_t -p tcp 9200-9300" >> $TMP

  /usr/sbin/semanage -S $selinuxvariant -i $TMP
fi
```

Deployment

[lzap@lzapx foreman-selinux]$ cat foreman-selinux-relabel

```
#!/bin/sh

/sbin/restorecon -ri $* /usr/share/foreman \
  /usr/share/katello \
  /var/lib/foreman \
  /var/run/foreman \
  /run/foreman \
  /var/log/foreman \
  /etc/foreman \
  /etc/puppet/node.rb \
  /etc/sysconfig/foreman* \
  /etc/rc.d/init.d/foreman* \
  /etc/logrotate.d/foreman* \
  /etc/cron.d/foreman*
```

Deployment

[lzap@lzapx foreman-selinux]$ cat foreman-selinux.spec | grep ...

```
%define selinux_variants targeted
%define selinux_modules foreman foreman-proxy


%build
# determine distribution name and version
%if 0%{?rhel} >= 6
%define distver rhel%{rhel}
%endif
%if 0%{?fedora} >= 18
%define distver fedora%{fedora}
%endif

# build policy
for selinuxvariant in %{selinux_variants}; do
    make clean all NAME=${selinuxvariant} DISTRO=%{distver} VERSION=%{version}
    for selinuxmodule in %{selinux_modules}; do
        mv ${selinuxmodule}.pp.bz2 ${selinuxmodule}-${selinuxvariant}.pp.bz2
    done
done
```

Deployment

```
%install
for selinuxvariant in %{selinux_variants}; do
    install -d %{buildroot}%{_datadir}/selinux/${selinuxvariant}
    for selinuxmodule in %{selinux_modules}; do
        install -p -m 644 ${selinuxmodule}-${selinuxvariant}.pp.bz2 \
          %{buildroot}%{_datadir}/selinux/${selinuxvariant}/${selinuxmodule}.pp.bz2
    done
done

make clean install-data NAME=${selinuxvariant} DISTRO=%{distver} \
                            VERSION=%{version} INSTPREFIX=%{buildroot}
```

Deployment

```
%post
if /usr/sbin/selinuxenabled; then
    # install and upgrade
    %{_sbindir}/%{name}-enable
fi

%posttrans
if /usr/sbin/selinuxenabled; then
    # install and upgrade
    %{_sbindir}/%{name}-relabel
fi

%preun
if /usr/sbin/selinuxenabled; then
    # uninstall only
    if [ $1 -eq 0 ]; then
        %{_sbindir}/%{name}-disable
    fi
    # upgrade and uninstall
    %{_sbindir}/%{name}-relabel
fi
```

Deployment

```
%files
%doc Contributors CHANGELOG LICENSE foreman.fc foreman.if foreman.te
%attr(0600,root,root) %{_datadir}/selinux/*/foreman.pp.bz2
%{_datadir}/selinux/devel/include/%{moduletype}/foreman.if
%attr(0755,root,root) %{_sbindir}/%{name}-enable
%attr(0755,root,root) %{_sbindir}/%{name}-disable
%attr(0755,root,root) %{_sbindir}/%{name}-relabel
%{_mandir}/man8/%{name}-enable.8.gz
%{_mandir}/man8/%{name}-disable.8.gz
%{_mandir}/man8/%{name}-relabel.8.gz
```

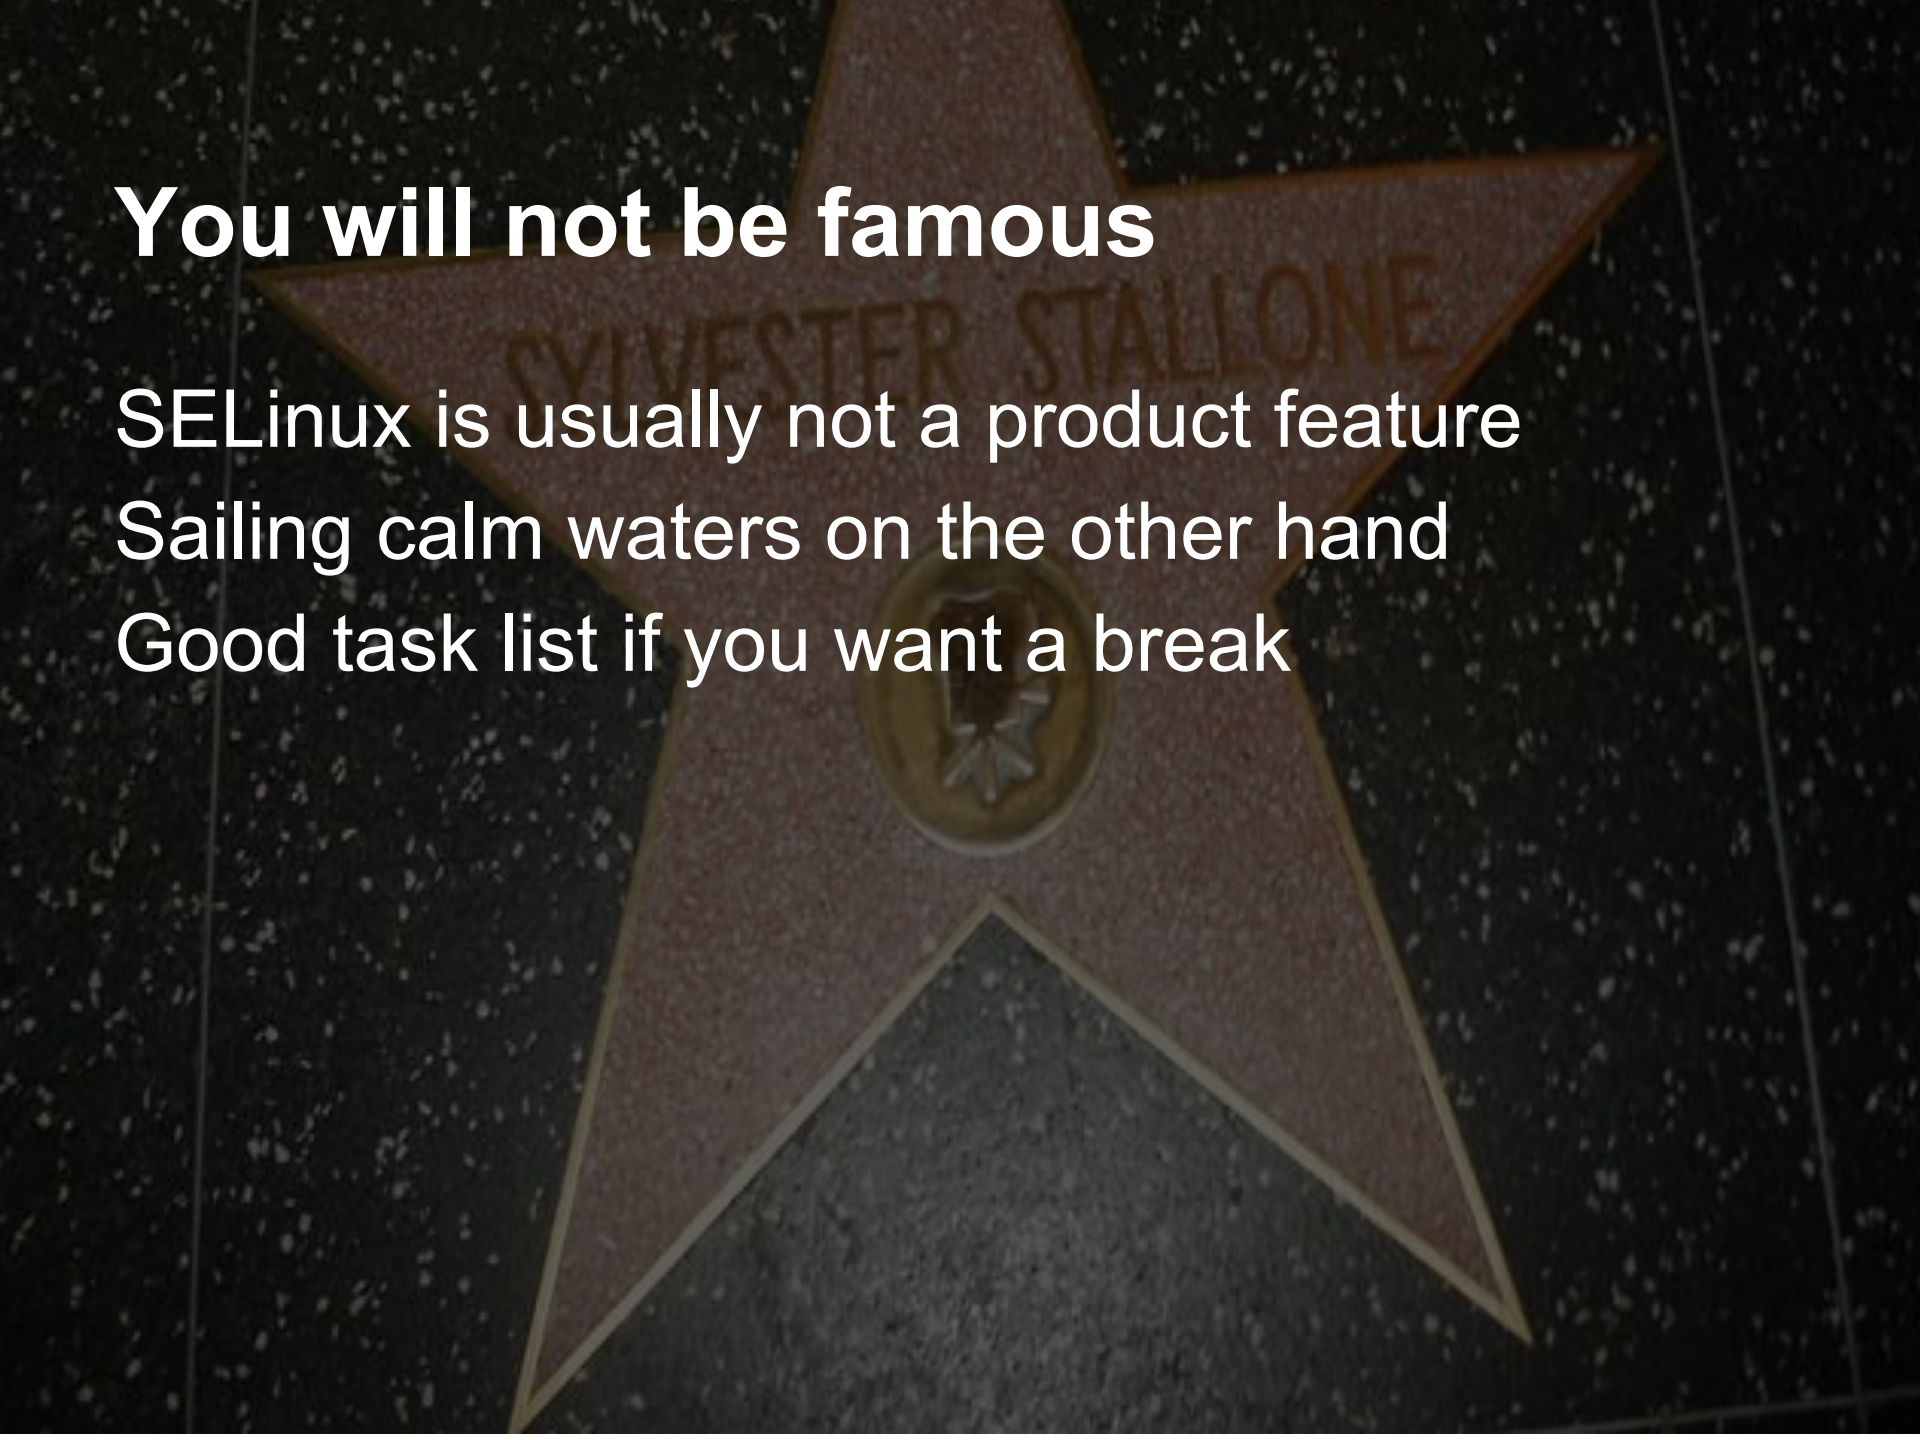

You will not be famous

SELinux is usually not a product feature

Sailing calm waters on the other hand

Good task list if you want a break

One more thing

How to file a SELinux bug

PROCESSES

```
ps axuZ
```

FILES

```
restorecon -rvn /
```

DENIALS

```
ausearch -m AVC
```


Q&_

Image credits - thanks

http://en.wikipedia.org/wiki/Joke_chess_problem#cite_note-1 (V. Ropke, Skakbladet 1942)

https://www.flickr.com/photos/x1brett/4600461689/

https://www.flickr.com/photos/nesster/3168425434/

https://openclipart.org/detail/4735/police-car-alarm-by-toplus

https://www.flickr.com/photos/caitlinator/3708011885/

http://aerokay.deviantart.com/art/The-Who-Poster-236014991

http://commons.wikimedia.org/wiki/File:PEO_M4_Carbine_RAS.jpg

http://commons.wikimedia.org/wiki/File:Horror_Images_Revolt_of_the_Zombies.jpg

http://commons.wikimedia.org/wiki/File:Blue_alarm_clock_(1).jpg

http://en.wikipedia.org/wiki/Big_Show

http://commons.wikimedia.org/wiki/File:INF_inspection.JPEG

http://pixabay.com/id/editor-teks-vim-perangkat-lunak-27620/

http://en.wikipedia.org/wiki/Smoking_in_Albania

http://commons.wikimedia.org/wiki/File:Bank-Security-Guard-Sleeping.jpeg

http://commons.wikimedia.org/wiki/File:Question_mark_(3534516458).jpg

https://openclipart.org/detail/182513/hazard-x-gold-by-Magirly-182513

http://pixabay.com/es/electricidad-flash-rayo-peligro-98819/

http://en.wikipedia.org/wiki/User:JustinTime55/sandbox/Apollo_11

http://en.wikipedia.org/wiki/Automotive_design